

No Hyperscaler? No Problem. Your Students also Break (and Fix) Networks

Constantine Doumanidis
Princeton University

Hongyu Hè
Princeton University

Maria Apostolaki
Princeton University

ABSTRACT

Network operators diagnose faults by issuing commands, forming hypotheses, and iterating to a fix. Yet, no dataset captures this diagnostic process. Production traces are locked down, and synthetic benchmarks reduce debugging to single injected faults on homogeneous topologies.

We propose mining students’ actions during the completion of networking course assignments as a source of realistic, labeled, and shareable debugging traces. We test this idea in the mini-Internet, an open-source platform in which students configure their part of a miniature of the Internet, naturally generating cascading, multi-protocol failures and the trajectories that fix them.

Realizing this opportunity requires confronting two challenges: what to log and at what granularity, and how to extract clean diagnostic trajectories from noisy, concurrent student sessions. We describe our initial design and the open research problems ahead. Our preliminary results show that non-trivial faults such as misconfigurations and diagnostic patterns can be mined from student logs automatically, and that the resulting guidance improves the root cause analysis accuracy of LLM agents.

1 INTRODUCTION

Managing networks is hard. Operators face a wide range of tasks: from routine work like adding a new server and ensuring traffic reaches it, to diagnosing configuration mistakes, to untangling cascading protocol failures whose root cause spans multiple systems and teams. The rise of LLMs has renewed interest in automating these tasks, but existing approaches still struggle: even on purpose-built troubleshooting benchmarks, today’s agents still struggle to localize faults and pinpoint root causes [26]. A key reason is data scarcity. Production environments generate rich operational traces, but they stay locked behind confidentiality walls. What is studied in research takes one of two forms: single injected faults on simulated homogeneous topologies, or sanitized industry snapshots scrubbed of most operational context. The messy, cascading, multi-party failures, the routine configuration (*i.e.*, the concrete steps an operator takes to configure a new router) and debugging work (*i.e.*, sequences of diagnostic actions) that define real network management are simply absent.

We argue that the networking research community is sitting on an overlooked gold mine: academic networking courses. Consider a course built around the mini-Internet [13], where each student group operates its own Autonomous System, initializes all routers’ configurations, establishes BGP sessions with neighboring ASes, and keeps their slice of the Internet alive all under the pressure of a semester deadline. Critically, these ASes are interconnected and share one server: a mistake in one AS becomes an outage in another, and a crashed container becomes someone else’s unexplained link failure. Thus, failures emerge organically, not according to a predetermined script, from the interaction of dozens of concurrent operators (other groups) making independent decisions: Group A’s route policy clashes with Group B’s, while Group C’s link goes down because their containers ran out of memory. Additionally, students perform fault diagnosis and mitigation, which often resembles genuine diagnostic reasoning similar to what we would expect from experienced operators. Students are, of course, not seasoned engineers, and we do not claim their faults or their reasoning match a production network run by one. What a course provides is a scarce ingredient: dozens of independent people with real stakes making concurrent decisions on a shared substrate, recorded. This cannot be emulated with, say, graduate students or engineers solving open-source challenges or injected faults. Critically, generated traces are also shareable: per-group topologies are identical, so no group is identifiable; IPs are abstracted; and any fault can be replayed on a fresh instance of the open-source platform.

Turning this observation into a research asset is not straightforward, and requires solving two families of problems that mirror open challenges in production monitoring. The first is instrumentation: deciding what to log, at what granularity, and how to do so without disrupting the students or the platform itself. The second is mining: extracting clean, meaningful patterns of faults or of diagnostic trajectories from sessions that are noisy, temporally ambiguous, and interleaved across concurrent operators. Two key insights drive our approach. To minimize load, we need to monitor what is strictly needed to reproduce students’ reasoning; hence, we only rarely run expensive global experiments such as PingMesh-like measurements but record each keystroke of students’ command lines. This allows us to capture not just the final network state, but the full operational process —

the sequence of commands, hypotheses, and configuration changes that led there. To make sense of recorded sessions, we transform them through a series of abstraction steps into a symbolic command sequence, and leverage carefully crafted templates to identify recurring patterns in the debugging processes found across student groups.

We instrumented one semester of our course: 78 ASes, 26 of them run by student groups, 26 devices each. Three templates alone mine cover 10090 occurrences (§4). The mined faults are not toy benchmarks: given full documentation and shell access to every device, a GPT-5.4-mini agent localizes our L2 subnetting fault in 3 of 5 runs, never localizes our BGP traffic engineering fault, and fails to even detect the latter in 4 of 5 runs. Runbooks generated automatically from the mined patterns lift root cause accuracy to 100% on both, with 32% and 38% fewer input tokens. Neither of our two problems is specific to the mini-Internet: what to log and how to mine it are the questions production monitoring has not answered either.

2 MOTIVATION

Next, we outline three concrete applications enabled by the data we collect from monitoring students as they work on the mini-Internet. While already useful in this setting, these applications reach their full potential only when the same methodology is applied to other populations (trained engineers, researchers) and other platforms, from educational testbeds to production networks.

Realistic, labeled, and shearable fault benchmark. The monitoring system produces a fault corpus that can serve as a reliable benchmark to develop, evaluate and compare RCA agents. Unlike recent troubleshooting benchmarks built from curated or injected incidents [26], it has properties no existing dataset offers. First, it contains real cascading failures caused by different individuals’ actions, involving various protocols and even hardware faults. For instance, Group A’s route policy interacts with Group B’s OSPF config, while Group C’s link fails because its containers crashed. Second, the dataset is reproducible and shareable: because we log the full config state and topology, any fault can be replayed on a fresh instance of the open-source mini-Internet platform. We make sure that no student or group is identifiable through their logs, as internal AS (per-group) topologies are identical and IPs are abstracted away. Irrelevant commands or things that users might have typed by accident are removed, and students were informed of our recording, which also reduced cheating attempts. Third, the dataset is naturally labeled — students eventually solve each problem, and the fix diff reveals the root cause without manual annotation.

Diagnosis Trajectories as demonstration data for network agents Diagnosis trajectories, *i.e.*, sequences of diagnostic commands, configuration changes, and their effect on connectivity, are useful for training network copilots, or serving as a realistic human baseline for evaluating RCA systems. Recent work reinforces their value: an agent that explicitly encodes the human, step-by-step diagnostic methodology substantially outperforms free-form reasoning at root-cause localization [22], suggesting the process itself, not just the eventual fix, carries useful signal. Yet such trajectories are not available today beyond locked-down production environments that cannot be shared. By recording students’ terminal sessions alongside config snapshots and ping mesh measurements, we obtain both successful trajectories (efficient paths from symptom to fix that serve as expert demonstrations) and failed ones (plausible-but-wrong hypotheses that were tried and ruled out, serving as negative examples). Together, they provide the reward signal needed for imitation learning or RLHF-style fine-tuning, with a built-in human baseline for evaluation.

Evidence-based Monitoring Design Network monitoring systems today are designed by intuition: operators add dashboards, metrics, and alerts until something feels sufficient, with no principled way to know whether they are collecting too much, too little, or the wrong things entirely. A data-driven approach would allow operators to identify redundant sources that can be dropped, and gaps where critical signals are missing. By aggregating which telemetry commands human operators reach for across hundreds of debugging episodes, and which actually help them, we get exactly this: if nobody ever checks a particular metric, it is likely dead weight; if two commands always appear together, one may be redundant; if a fault type consistently requires information that no standard command surfaces, that is a monitoring gap. Of course, student decisions and preferences might not reflect the true monitoring value perfectly, yet it might reveal gaps in their understanding. More importantly, we believe this is a novel applications that a logging system like the one we propose can enable.

Designing better network courses In short: We can mine common error patterns that imply certain concepts need better instruction / reinforcement by instructors. We can also study how and when students employ the use of agents in their workflows, and understand what worked and what didn’t.

3 INITIAL DESIGN

This section describes the challenges we encountered and the insights that influenced the design of our data collection and interpretation methods. We stress that these hold beyond the mini-Internet and even the classroom setting.




Figure 1: Student command sequences are transformed through a series of abstraction steps to abstract sequences that expose common patterns among individual command sequences. The corpus of abstract sequences is then matched against a set of high-level templates (e.g., write after write with some commands in between). Each template can expose multiple individual patterns by constraining some of their parameters. Mined patterns can point to common faults such as configurationisms and/or guide LLMs in RCA, or routine tasks.

3.1 Instrumentation Challenges & Insights

Scale logging by avoiding redundancy: The full state of our mini-Internet network can be fully captured by recording extensive network data (e.g., link statistics, ARP tables, BGP databases, switch configuration files). However, such comprehensive data collection would be unscalable; as it would compete with students for system resources. We observe that there is a lot of redundancy in the data one would collect as, some signals can be derived by re-running the mini-Internet. For instance, we frequently log configuration files of all routers, but we only infrequently collect global data-plane measurements such as pings and traceroutes.

Capture Students’ Reasoning by recording terminals: Our goal is to study and understand the thought processes of students as they manage their networks. One way to do this is by asking students to keep detailed logs of their actions (config files, debug efforts, etc.). However, students might omit their failed attempts, or unnecessary actions, or might simply forget to do so. Our insight is that students’ thought processes manifest in their terminal sessions, which can be logged via commodity terminal recorders, like asciinema [18]. We augment asciinema for our use by making it tag terminal lines with precise timestamps. Since students can have concurrent active terminal sessions, timestamping allows us to merge their actions into a single timeline. To ensure the completeness of our terminal logs, we modify asciinema to write out the full primary virtual screen buffer rather than just what is displayed on the screen. We also later reduce noise in our analysis, by discarding sessions that were either uneventful (e.g., login-logout) or failed to properly start (e.g., due to exhausted system resources). At the end, our pipeline yields a sequence of user terminal commands and their corresponding outputs.

Exploit (not avoid) Failures: Despite our best efforts to maintain high availability for mini-Internet containers running on our server, we occasionally observed that a small

number of them crashed unexpectedly and needed to be re-stored. Rather than upgrading the server to eliminate these failures, we treated them as an opportunity to introduce random faults like those that commonly occur in real-world environments. Students were aware that such crashes could occur and were instructed on how to diagnose them and have their affected containers restored.

3.2 Interpretation Challenges & Insights:

Action	Description	Example Commands
create	Initialize resources	ip addr add ..., ovs-vsctl add-port...
read	Inspect resources	show ip bgp, do show running-config
update	Modify resources	ovs-vsctl set ..., restore_configs.sh
delete	Remove resources	ip addr del ..., no neighbor ...
probe	Probe network	ping 1.1.1.1, traceroute 1.1.1.1
navigate	Shell context changes	ssh root@ATLrouter, config terminal
system	Misc. commands	ls -l, cd, cat, chmod, save_configs.sh

Table 1: Actions Taxonomy: abstracting commands by the effect they have, facilitates mining patterns.

Resource	Description	Example Commands
interface	Links, switch ports & tunnels	ip tunnel add..., ifconfig eth0,
bridge	OVS Bridges	ovs-vsctl set-fail-mode br0 secure
address	IPs, MAC addrs.	ip addr add..., ip neighbor show
protocol	Routing config. & static routes	neighbor..., router bgp 100
policy	Filter rules, ACLs & route maps	route-map MAP permit 10
state	Status checks, diagnostic probes	ping 1.1.1.1, show ip bgp summary
config	Static config files	restore_configs.sh
system	Misc system operations.	ls -l, uname

Table 2: Resource Type Taxonomy: this abstraction enables mining patterns by relating commands through the network resources they access or modify.

Expose patterns by abstracting command sequences: Mining for patterns on raw commands is challenging since it is unlikely that any non-trivial sequence of commands appears verbatim more than a handful of times. Command strings can contain environment-specific identifiers (e.g., IP addresses, interface numbers), reflect specific user preferences (e.g., output formatting flags), and can have multiple

other syntactically different yet functionally equivalent alternatives from which users can select. Consider a simple host configuration pattern that we would like to spot: a student adds an IP address to an interface, they then set a default gateway, then proceed to check if both of these settings were applied to the system, and finally check connectivity to the upstream router. Student A may choose to configure their 55_A_EUH host using `ipaddr add 55.200.0.1/24 dev 55-S2`, followed by `ip route add default via 55.102.0.2`, then `ipaddr show` and `ip r`, and finally `ping 55.102.0.2`. Student B takes a different approach on their 24_A_EUH host: `ip address add 24.200.0.11/24 dev 24-S2`, `ifconfig`, then `ip route add default via 24.200.0.254`, `ip route list default`, and finally `ping group24.ATLrouter`. Both sequences are underpinned by the same basic host configuration and verification workflow, yet their command sequences appear very different in their current form.

To extract our desired pattern, we must first lift this sequence to a higher level of abstraction. Concretely, we identify three necessary transformations: First, we map domain-specific identifiers, such as IP addresses and AS-specific hostnames, to symbolic counterparts that indicate their positions in the topology. We are able to do this because each student AS in the mini-Internet has the exact same internal structure. For instance, in the above example we map both 55.102.0.2 and group24.ATLrouter to ATLrouter, 55-S2 and 24-S2 become S2, and so on. We next transform all commands that have multiple (short-cut type) alternative forms and notations to a standard form; in our example, we overwrite `ipaddr add` with `ip address add`, and `ip r` with `ip route`. Finally, we map commands to equivalence classes that we define using (action, resource) tuples. Our actions taxonomy, shown in Table 1, is inspired by CRUD *i.e.*, the Create, Read, Update, and Delete actions from database systems [19]. In our context, these actions can be used to initialize, view, modify, and remove network-related configurations. We augment our actions to include probe, which represents network measurements through reachability tests and traces. To capture students moving between different shell environments (*e.g.*, SSH, privileged configuration modes), we defined the navigate action. Finally, to categorize commands unrelated to network management, we created a system action. Similar in spirit, our resource taxonomy in Table 2 represents the entities that students interact with at a high level in their assignment, including interfaces, addresses, protocols, policies, static configurations, etc. In our running example, for instance, `ipaddr show` and `ifconfig` are both mapped to (read, address).

Mine patterns through carefully crafted templates: Identifying command patterns that are common is challenging because the potential candidates are too many to enumerate. Observe that all sequences of any size, composed of any type

of command, and even their permutations, should be checked to determine if they are common. Straightforward splitting the command sequences in time to reduce the search space *e.g.*, by session does not work: students often use multiple sessions concurrently, or might not conclude their actions within one session.

Instead, we propose mining for patterns in two steps, each partitioning command sequences along different axes. We start by extracting sequences that match high-level *templates* which aim to frame a semantically meaningful job that a student completed, such as fixing a misconfiguration, effectively reducing its size. Hence, templates are composed of interleaved sequences of *anchor* points and *unordered sets* of commands. Anchors aim to identify a change in state (*e.g.*, through a write-config change) or a trigger to action (*e.g.*, a through a read-failed ping). Once sequences are partitioned and grouped under a given template, we proceed to the second step, where we identify concrete patterns by specifying more variables of a given template. For example, we can pin down device, action, and resource type. Through this process, template #2 shown in Table 3 eventually leads to mining the patterns shown in 4. To illustrate the usefulness of templates, recall in our running example that Student A ran their IP and gateway verification commands after their configurations, whereas Student B interleaved theirs. Because template #1 in Table 3 starts with an unordered set, it can capture the configuration pattern of both students for this host.

Clean-up mined patterns via selective re-runs: Mined patterns identify commands that frequently occur together across our recordings. However, not all commands within a pattern contribute to the observed outcome. Furthermore, because of challenges related to recording and processing student terminal sessions, our command records may be incomplete. Without additional reasoning, it is hard to determine statically whether a mined pattern is *sufficient* to produce the observed outcome, and which of its commands are *necessary* for doing so. We overcome this by leveraging the reproducibility of the mini-Internet. For instance, in our running example we wish to determine which of the write commands in the pattern we mined are *sufficient and necessary* for A_EUH to ping host A_CHA which is connected to it via the same switch and has an IP address in the same /24 prefix. We restore the network state to the point in time before the student configured A_EUH, and evaluate every combination of the two create commands, followed by a ping from A_EUH to A_CHA and an environment reset. After these experiments, we can conclude that configuring only A_EUH’s interface with an IP address was sufficient for it to successfully ping A_CHA, and that the default gateway configuration was not necessary for this objective.

#	Template	# Occurrences
1	$S[d:d', a: "(create) (read) (update)", r: *]\{10\} \leftarrow A[d:d', a: "probe", r: *, probe.success: True, order(probe): 1]$	2087
2	$A[d:d', a: "(create) (update)", r: r'] - S[d:d', a: "(read) (probe)", r: *]\{10\} \leftarrow A[d:d', a: "(create) (update)", r: r']$	3678
3	$A[d: *, a: "probe", r: *, probe.success: False] \rightarrow S[d: *, a: "(create) (update) (delete)", *, *]\{10\}$	4325

Table 3: Example templates and number of occurrences. Notation: A - anchor command, S - unordered commands set. $\leftarrow$ / $\rightarrow$ - search order for S commands w.r.t the attached A. A, and S are parameterized by [d(evice), a(ction), r(esource), other args.]. S are also parameterized by their maximum size noted in {}.

4 INITIAL PROMISE

We demonstrate the feasibility of our method and showcase the potential of the patterns we mined from our class’ mini-Internet environment. Our setup comprised of 78 virtual ASes, 26 of which were managed by student groups of 2 or 3 and contained identical topologies of 26 devices (e.g., routers, switches, hosts).

Claim #1: Patterns exist and can be mined. Applying our pattern mining method to traces collected from the mini-Internet yields hundreds of recurring patterns, demonstrating that student behavior contains mineable structure. We illustrate this using the three example templates shown in Table 3. Template #1 targets *device configuration* patterns such as that of the example in Section 3. It is comprised of an unordered set of up to 10 commands that read or modify network state on a given device, followed by the first successful probe command executed on the same device. This template occurs 2087 times in our dataset. Template #2 mines *misconfiguration diagnosis* patterns: an anchor that modifies a specific resource on a given device is followed by commands that probe or read the network state, and are finally followed by an anchor command that overwrites the configuration of the same resource on the same device. This template manifests 3678 times in total in our data. Template #3 targets patterns that fix *connectivity issues*: a failed probe anchor is immediately followed by a set of commands that modify network state. This template occurs 4325 times across the dataset. Across these three templates alone, our method discovers more than 1,500 unique patterns, indicating that student terminal sessions exhibit structure that can be mined automatically.

Claim #2: Mined patterns are non-trivial RCA benchmarks. We evaluate two of our mined misconfiguration patterns as RCA benchmarks using a GPT-5.4 mini-based RCA agent. Despite being provided with detailed network documentation and full access to the topology, the agent struggles to identify the root causes of the faults. Our L2 Subnetting misconfiguration captures an error where students configure two router subinterfaces with overlapping IP prefixes, causing the traffic to one of the downstream VLANs to be dropped. Our BGP traffic engineering misconfiguration requires configuring AS_PATH prepending to correctly steer inbound traffic. As shown in Table 5, the agent successfully localized the L2 Subnetting misconfigurations in only 60% of

runs, and failed to do so for the BGP traffic engineering misconfiguration. The RCA success rates were likewise 60% and 0% respectively. These failures occurred despite substantial efforts: across 5 diagnosis runs for each misconfiguration, the agent took 10 steps, made more than 15 tool invocations, and consumed over 69,000 and 91,000 input tokens for the L2 subnetting and BGP traffic engineering misconfigurations respectively.

Claim #3: Mined runbooks can guide RCA agents. We augment our RCA agent with access to runbooks that we automatically generated from the patterns corresponding to the mined misconfigurations and re-evaluate its performance. In this setting, the runbooks substantially improve the agent’s performance (Table 5). The agent now achieves high RCA success rate on both L2 subnetting and BGP traffic engineering. The runbooks also reduce the effort required for diagnosis: the L2 subnetting misconfiguration required approximately 6 steps and 12 tool calls, while consuming 32% fewer tokens. For the BGP traffic engineering misconfiguration, the agent required more than a step less, just over 11 tool calls and 38% fewer input tokens. These results suggest that mined patterns can be transformed into useful troubleshooting guidance for RCA agents.

5 RELATED WORK

Network troubleshooting and AI benchmarks: A long line of work automates fault diagnosis and management. Early systems localize faults from production telemetry, including NetPoirot from endpoint TCP statistics [2] and 007 from end-host signals [1]. Recent systems drive operations with LLMs, spanning failure localization in production networks [24], dialogue-based diagnosis [25], intent-driven management at hyperscale [27], code generation for management tasks [16], and a vision for AI-driven incident management [11]. To measure progress, shared benchmarks such as NIKA replay curated troubleshooting incidents [26], and NetConfEval scores configuration synthesis from natural-language intent [23]. These systems learn from proprietary production incidents that are never released, and the benchmarks rely on faults injected or curated by their authors. Neither pairs a fault with the commands, hypotheses, and dead ends that a human followed to fix it, which is the signal we set out to capture.

Misconfiguration	Pattern	# Occurrences
L2 Subnetting	[d:"ATLrouter", a:"create", r:"address"] - [[d:"a", a:"probe", r:"state"], [d:"a", a:"read", r:"state"]] - [d:"ATLrouter", a:"delete", r:"address"]	41
BGP TE	[d:"MSProuter", a:"update", r:"protocol"] - [[d:"a", a:"probe", r:"state"], [d:"a", a:"read", r:"state"]] - [d:"MSProuter", a:"update", r:"protocol"]	89

Table 4: Patterns used to mine the L2 subnetting and BGP traffic engineering misconfigurations.

Misconfiguration	Runbooks	Steps	# Tools	# In Tokens	# Out Tokens	DR(%)	LR(%)	RCA(%)
L2 Subnetting	False	10.0	15.4	69302.8	890.4	100.0%	60.0%	60.0%
BGP TE	False	10.6	15.6	91566.2	775.6	20.0%	0.0%	0.0%
L2 Subnetting	True	5.8	11.6	47096.8	664.2	100.0%	100.0%	100.0%
BGP TE	True	9.2	11.2	56882.2	611.4	100.0%	80.0%	100.0%

Table 5: Mined misconfigurations are challenging benchmarks for RCA agents. Providing the agent with access to runbooks automatically generated from mined patterns significantly improves performance. Notation: DR - detection success rate, LR - localization success rate, RCA - root cause analysis success rate.

Educational and emulation platforms: Reproducible environments for building and running networks are mature. Mininet prototypes single-machine SDNs [15], Emulab and GENI provision controlled experiments at scale [5, 28], and Katharà assembles container-based topologies [6]. The mini-Internet brings the same reproducibility to teaching, letting student groups jointly operate a miniature Internet [13]. We repurpose such a platform into an instrument that records and mines how its operators reason. The platforms supply the substrate, but none captures the human debugging process running on top of it.

Operational datasets and operator studies: Understanding real networks has relied on production data held by a single organization. Telemetry systems such as Pingmesh continuously measure large data centers [10], and measurement studies characterize why networks fail [9, 20]. Related work quantifies configuration and management complexity [4] and analyzes the data plane implied by configurations [8]. Field studies instead observe how administrators work in practice [3]. Each source is locked to one provider, summarizes outcomes after the fact rather than the diagnostic trajectory, or reports qualitative findings that yield no labeled, minable traces.

Mining the human process: Our approach builds on educational data mining, where capturing the process rather than the final artifact has proven its worth. Repositories such as Blackbox record millions of novice programming sessions [7], compilation behavior has been mined across tens of thousands of users [14], and program trajectories drive automated feedback [21]. We carry the same principle into network operations, where the debugging trajectory is the asset worth keeping. Sequential pattern mining offers a natural starting point for finding recurring structure [12, 17]. Standard methods assume a single ordered stream, however, and cannot cleanly separate the interleaved, concurrent sessions our setting produces, which motivates the template-based mining of Section 3.

6 FUTURE AGENDA

What happens once the mini-Internet structure is gone?

Our abstraction leans on two properties of the mini-Internet: (i) every group configures an AS with the same internal topology, so a name resolves to a position and a position to a role; (ii) every group is set the same assignments, so equivalent tasks recur and give our templates a frequency signal. Yet, a production network might not be symmetrical enough to resolve names to roles, and tasks in production are way more diverse. Finding similar abstractions is hence an open question. Roles might come from text: textbooks, vendor documentation, and RFCs already describe networks the way we want to abstract them, naming roles rather than devices and decomposing tasks into canonical subtasks. Significance through frequency might need to be replaced by significance through authority *e.g.*, an observed sequence by an expert might be a pattern or a candidate pattern without further support. Further focusing on patterns anchored in data plane signals or observable effects should be a promising direction as they are transferable: a newly failing probe means the same thing whatever configured it.

Networking courses as research infrastructure: Our work suggests that educational networking platforms can serve a second purpose beyond instruction: producing realistic, reproducible, and readily shareable operational datasets. This raises a broader question for our community: how can we co-design networking courses to simultaneously optimize student learning and the generation of high-fidelity research artifacts? By introducing standardized topologies, reproducible assignments, and privacy-preserving instrumentation multiple institutions could collaboratively build corpora of configuration tasks, debugging trajectories, and failures cannot be obtained from production networks. Such datasets would benefit not only research in AI-driven network management, but also networking education by revealing common misconceptions and gaps in instruction.

REFERENCES

- [1] Behnaz Arzani, Selim Ciraci, Luiz Chamon, Yibo Zhu, Hongqiang Harry Liu, Jitu Padhye, Boon Thau Loo, and Geoff Outhred. 2018. 007: Democratically Finding the Cause of Packet Drops. In *15th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*. USENIX Association, 419–435. <https://www.usenix.org/conference/nsdi18/presentation/arzani>
- [2] Behnaz Arzani, Selim Ciraci, Boon Thau Loo, Assaf Schuster, and Geoff Outhred. 2016. Taking the Blame Game out of Data Centers Operations with NetPoirot. In *Proceedings of the 2016 ACM SIGCOMM Conference*. Association for Computing Machinery, 440–453. <https://doi.org/10.1145/2934872.2934884>
- [3] Rob Barrett, Eser Kandogan, Paul P. Maglio, Eben M. Haber, Leila A. Takayama, and Madhu Prabaker. 2004. Field Studies of Computer System Administrators: Analysis of System Management Tools and Practices. In *Proceedings of the 2004 ACM Conference on Computer Supported Cooperative Work (CSCW)*. Association for Computing Machinery, 388–395. <https://doi.org/10.1145/1031607.1031672>
- [4] Theophilus Benson, Aditya Akella, and David A. Maltz. 2009. Unraveling the Complexity of Network Management. In *6th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*. USENIX Association, 335–348. <https://www.usenix.org/conference/nsdi-09/unraveling-complexity-network-management>
- [5] Mark Berman, Jeffrey S. Chase, Lawrence Landweber, Akihiro Nakao, Max Ott, Dipankar Raychaudhuri, Robert Ricci, and Ivan Seskar. 2014. GENI: A Federated Testbed for Innovative Network Experiments. *Computer Networks* 61 (2014), 5–23. <https://doi.org/10.1016/j.bjp.2013.12.037>
- [6] Gaetano Bonofiglio, Veronica Iovinella, Gabriele Lospoto, and Giuseppe Di Battista. 2018. Katharà: A container-based framework for implementing network function virtualization and software defined networks. In *NOMS 2018 - 2018 IEEE/IFIP Network Operations and Management Symposium*. 1–9. <https://doi.org/10.1109/NOMS.2018.8406267>
- [7] Neil C. C. Brown, Michael Kölling, Davin McCall, and Ian Utting. 2014. Blackbox: A Large Scale Repository of Novice Programmers’ Activity. In *Proceedings of the 45th ACM Technical Symposium on Computer Science Education (SIGCSE)*. Association for Computing Machinery, 223–228. <https://doi.org/10.1145/2538862.2538924>
- [8] Ari Fogel, Stanley Fung, Luis Pedrosa, Meg Walraed-Sullivan, Ramesh Govindan, Ratul Mahajan, and Todd Millstein. 2015. A General Approach to Network Configuration Analysis. In *12th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*. USENIX Association, 469–483. <https://www.usenix.org/conference/nsdi15/technical-sessions/presentation/fogel>
- [9] Phillipa Gill, Navendu Jain, and Nachiappan Nagappan. 2011. Understanding Network Failures in Data Centers: Measurement, Analysis, and Implications. In *Proceedings of the ACM SIGCOMM 2011 Conference*. Association for Computing Machinery, 350–361. <https://doi.org/10.1145/2018436.2018477>
- [10] Chuanxiong Guo, Lihua Yuan, Dong Xiang, Yingnong Dang, Ray Huang, Dave Maltz, Zhaoyi Liu, Vin Wang, Bin Pang, Hua Chen, Zhi-Wei Lin, and Varugis Kurien. 2015. Pingmesh: A Large-Scale System for Data Center Network Latency Measurement and Analysis. In *Proceedings of the 2015 ACM SIGCOMM Conference*. Association for Computing Machinery, 139–152. <https://doi.org/10.1145/2785956.2787496>
- [11] Pouya Hamadanian, Behnaz Arzani, Sadjad Fouladi, Siva Kesava Reddy Kakarla, Rodrigo Fonseca, Denizcan Billor, Ahmad Cheema, Edet Nkposong, and Ranveer Chandra. 2023. A Holistic View of AI-driven Network Incident Management. In *Proceedings of the 22nd ACM Workshop on Hot Topics in Networks (HotNets)*. Association for Computing Machinery, 180–188. <https://doi.org/10.1145/3626111.3628176>
- [12] Jiawei Han, Jian Pei, Behzad Mortazavi-Asl, Helen Pinto, Qiming Chen, Umeshwar Dayal, Meichun Hsu, et al. 2001. Prefixspan: Mining sequential patterns efficiently by prefix-projected pattern growth. In *proceedings of the 17th international conference on data engineering*. IEEE Piscataway, NJ, USA, 215–224.
- [13] Thomas Holterbach, Tobias Bühler, Tino Rellstab, and Laurent Vanbever. 2020. An Open Platform to Teach How the Internet Practically Works. *SIGCOMM Comput. Commun. Rev.* (2020). <https://doi.org/10.1145/3402413.3402420>
- [14] Matthew C. Jadud and Brian Dorn. 2015. Aggregate Compilation Behavior: Findings and Implications from 27,698 Users. In *Proceedings of the Eleventh Annual International Conference on International Computing Education Research (ICER)*. Association for Computing Machinery, 131–139. <https://doi.org/10.1145/2787622.2787718>
- [15] Bob Lantz, Brandon Heller, and Nick McKeown. 2010. A Network in a Laptop: Rapid Prototyping for Software-Defined Networks. In *Proceedings of the 9th ACM SIGCOMM Workshop on Hot Topics in Networks (HotNets)*. Association for Computing Machinery, Article 19, 6 pages. <https://doi.org/10.1145/1868447.1868466>
- [16] Sathiyaa Kumaran Mani, Yajie Zhou, Kevin Hsieh, Santiago Segarra, Trevor Eberl, Eliran Azulai, Ido Frizler, Ranveer Chandra, and Srikanth Kandula. 2023. Enhancing Network Management Using Code Generated by Large Language Models. In *Proceedings of the 22nd ACM Workshop on Hot Topics in Networks (HotNets)*. Association for Computing Machinery, 196–204. <https://doi.org/10.1145/3626111.3628183>
- [17] Heikki Mannila, Hannu Toivonen, and A. Inkeri Verkamo. 1995. Discovering Frequent Episodes in Sequences. In *Proceedings of the First International Conference on Knowledge Discovery and Data Mining (KDD)*. AAAI Press, 210–215. <https://aaai.org/papers/kdd95-024-discovering-frequent-episodes-in-sequences/>
- [18] Kulik Marcin. 2026. asciinema. <https://github.com/asciinema/asciinema>
- [19] J. Martin. 1983. *Managing the Data-base Environment*. Prentice-Hall. <https://books.google.com/books?id=nMgmAAAAMAAJ>
- [20] Justin Meza, Tianyin Xu, Kaushik Veeraraghavan, and Onur Mutlu. 2018. A Large Scale Study of Data Center Network Reliability. In *Proceedings of the Internet Measurement Conference (IMC)*. Association for Computing Machinery, 393–407. <https://doi.org/10.1145/3278532.3278566>
- [21] Chris Piech, Jonathan Huang, Andy Nguyen, Mike Phulsuksombati, Mehran Sahami, and Leonidas Guibas. 2015. Learning Program Embeddings to Propagate Feedback on Student Code. In *Proceedings of the 32nd International Conference on Machine Learning (ICML) (PMLR, Vol. 37)*. PMLR, 1093–1102. <https://proceedings.mlr.press/v37/piech15.html>
- [22] Kuan-Hao Tseng, Niruth Bogahawatta, Yasod Ginige, Kosta Dekic, Arunan Sivanathan, and Suranga Seneviratne. 2026. SADE: Symptom-Aware Diagnostic Escalation for LLM-Based Network Troubleshooting. arXiv:2605.04530 [cs.NI] <https://arxiv.org/abs/2605.04530>
- [23] Changjie Wang, Mariano Scazzariello, Alireza Farshin, Simone Ferlin, Dejan Kostić, and Marco Chiesa. 2024. NetConfEval: Can LLMs Facilitate Network Configuration? *Proceedings of the ACM on Networking* 2, CoNEXT2, Article 7 (2024), 25 pages. <https://doi.org/10.1145/3656296>
- [24] Chenxu Wang, Xumiao Zhang, Runwei Lu, Xianshang Lin, Xuan Zeng, Xinlei Zhang, Zhe An, Gongwei Wu, Jiaqi Gao, Chen Tian, Guihai Chen, Guyue Liu, Yuhong Liao, Tao Lin, Dennis Cai, and Ennan Zhai. 2025. Towards LLM-Based Failure Localization in Production-Scale Networks. In *Proceedings of the ACM SIGCOMM 2025 Conference* (São Francisco Convent, Coimbra, Portugal) (SIGCOMM ’25). Association for Computing Machinery, New York, NY, USA, 496–511. <https://doi.org/10.1145/3718958.3750505>

- [25] Haopei Wang, Anubhavnidhi Abhashkumar, Changyu Lin, Tianrong Zhang, Xiaoming Gu, Ning Ma, Chang Wu, Songlin Liu, Wei Zhou, Yongbin Dong, Weirong Jiang, and Yi Wang. 2024. NetAssistant: Dialogue Based Network Diagnosis in Data Center Networks. In *21st USENIX Symposium on Networked Systems Design and Implementation (NSDI 24)*. USENIX Association, Santa Clara, CA, 2011–2024. <https://www.usenix.org/conference/nsdi24/presentation/wang-haopei>
- [26] Zhihao Wang, Alessandro Cornacchia, Alessio Sacco, Franco Galante, Marco Canini, and Dingde Jiang. 2025. A Network Arena for Benchmarking AI Agents on Network Troubleshooting. arXiv:2512.16381 [cs.NI] <https://arxiv.org/abs/2512.16381>
- [27] Zhaodong Wang, Samuel Lin, Guanqing Yan, Soudeh Ghorbani, Minlan Yu, Jiawei Zhou, Nathan Hu, Lopa Baruah, Sam Peters, Srikanth Kamath, Jerry Yang, and Ying Zhang. 2025. Intent-Driven Network Management with Multi-Agent LLMs: The Confucius Framework. In *Proceedings of the ACM SIGCOMM 2025 Conference*. Association for Computing Machinery, 347–362. <https://doi.org/10.1145/3718958.3750537>
- [28] Brian White, Jay Lepreau, Leigh Stoller, Robert Ricci, Shashi Guruprasad, Mac Newbold, Mike Hibler, Chad Barb, and Abhijeet Joglekar. 2002. An Integrated Experimental Environment for Distributed Systems and Networks. In *5th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*. USENIX Association, 255–270. <https://www.usenix.org/conference/osdi-02/integrated-experimental-environment-distributed-systems-and-networks>